

RAG Knowledge Base: Deep Neural Networks (NN) - Single File

ID: RN-01 | Topic: Fundamentals | Question (P): What is an **Artificial Neuron** (Perceptron)? | **Answer (R):** It's the basic computational unit. It receives inputs (x), multiplies them by **weights** (W), adds a **bias** (b), and passes the result through an **activation function** (g) to produce an output.

ID: RN-02 | Topic: Fundamentals | Question (P): What is the difference between a *Shallow* Neural Network and a **Deep Neural Network (DNN)**? | **Answer (R):** A *shallow* network has only one or two hidden layers. A **DNN** has **multiple hidden layers** (generally three or more), which allows it to learn hierarchical representations and complex features from data.

ID: RN-03 | Topic: Architecture | Question (P): What are the **Weights** (W) and **Bias** (b) in an NN? | **Answer (R):** The **Weight** is the parameter that determines the importance of an input connection. The **Bias** is a constant value added to the weighted sum to shift the activation threshold. **These are the parameters the network learns.**

ID: RN-04 | Topic: Architecture | Question (P): What is the key purpose of the **Activation Function** (e.g., ReLU, Sigmoid)? | **Answer (R):** It introduces **non-linearity** to the network. Without it, the network would only calculate linear combinations of the inputs, limiting its capacity to model complex relationships.

ID: RN-05 | Topic: Architecture | Question (P): What is a **Hyperparameter** (e.g., Learning Rate, number of layers)? | **Answer (R):** It is a parameter whose value is **set manually** before the training process begins. It is distinct from weights and biases, which are parameters the network learns automatically.

ID: RN-06 | Topic: Architecture | Question (P): What does the **Softmax** function do in the output layer? | **Answer (R):** It is commonly used in **multiclass classification** problems. It converts the output of the final layer into a **probability distribution**, where the sum of all outputs is 1, indicating the likelihood that the input belongs to each class.

ID: RN-07 | Topic: Feedforward | Question (P): What is the **Feedforward** process (Forward Propagation)? | **Answer (R):** It is the **prediction calculation** process. Input data (X) travels through the network, layer by layer, using the current weights and biases to generate the **final output or prediction** (y^{\wedge}).

ID: RN-08 | Topic: Feedforward | Question (P): How is the **net activation** (Z) of a layer represented in terms of matrices? | **Answer (R):** It is represented as a matrix multiplication for the entire *minibatch* or layer: $Z = W \cdot A_{prev} + B$. This allows all neurons to be processed efficiently, which is critical for GPU performance.

ID: RN-09 | Topic: Feedforward | Question (P): What is the mathematical formula for a neuron's activation (a)? | **Answer (R):** The neuron's output is calculated as: $a=g(z)$, where $z=(\sum_i w_{ix}i)+b$. The term z is called the **weighted sum** or **net activation**.

ID: RN-10 | Topic: Functions | Question (P): Why is the **ReLU** function the most widely used in hidden layers? | **Answer (R):** The Rectified Linear Unit (ReLU) is simple to compute ($\max(0,z)$) and, crucially, it helps **mitigate the vanishing gradient problem** compared to Sigmoid or Tanh, especially in deep networks.

ID: RN-11 | Topic: Learning | Question (P): What is the **Cost Function** (or Loss Function)? | **Answer (R):** It is a metric that quantifies the network's **error** by measuring the difference between the prediction (\hat{y}) and the actual value (y). The goal of training is to find the weights that **minimize** this function.

ID: RN-12 | Topic: Backpropagation | Question (P): What is the **Backpropagation** process? | **Answer (R):** It is the algorithm that efficiently computes the **gradients** (the partial derivative of the cost function with respect to every parameter, $\partial W \partial L$) to determine how much each weight contributed to the network's total error.

ID: RN-13 | Topic: Backpropagation | Question (P): What is the role of the **Chain Rule** in *Backpropagation*? | **Answer (R):** The chain rule from calculus is essential. It allows the total gradient calculation to be **decomposed** into the product of local gradients for each layer and neuron, propagating the error **backward** through the network.

ID: RN-14 | Topic: Optimization | Question (P): What is **Gradient Descent** and what does it do? | **Answer (R):** It is the optimization algorithm that uses the **gradients** calculated by *Backpropagation* to **iteratively update the weights and biases** in the direction opposite to the gradient (the direction of **steepest descent**) to minimize the cost function.

ID: RN-15 | Topic: Optimization | Question (P): What is the **Learning Rate (α)**? | **Answer (R):** It is a key **hyperparameter** that defines the **step size** in each iteration of Gradient Descent. A high rate can "overshoot" the minimum, while a low rate makes training very slow.

ID: RN-16 | Topic: Matrices/Gradients | Question (P): How is the calculation of the **weight gradient ($\partial W \partial L$)** represented in matrices? | **Answer (R):** The gradient is calculated as a matrix product (typically $A_{prev}^T \cdot \delta$), combining the error vector of the current layer (δ) with the activation of the previous layer (A_{prev}) to obtain the weight updates for the entire layer.

ID: RN-17 | Topic: Optimization | Question (P): What is the difference between **Stochastic Gradient Descent (SGD)** and *Batch* Gradient Descent? | **Answer (R):** **SGD** updates weights after every **single data sample**, introducing more noise but potentially finding better

minima. *Batch* Gradient Descent updates weights after processing the **entire training set**, which is slower but more stable.

ID: RN-18 | Topic: Optimization | Question (P): How does an optimizer like **ADAM** work? |

Answer (R): ADAM (Adaptive Moment Estimation) is an optimizer that **adaptively adjusts the learning rate** for each weight individually, using an exponential moving average of both the gradients (momentum) and the squared gradients (velocity) to accelerate convergence.

ID: RN-19 | Topic: Regularization | Question (P): What are **Overfitting** and **Underfitting**? |

Answer (R): **Overfitting** occurs when the network learns the noise in the training data too well, failing to generalize. **Underfitting** occurs when the network is too simple and fails to capture the essential relationships in the data.

ID: RN-20 | Topic: Regularization | Question (P): How does the **Dropout** technique help prevent *Overfitting*? | **Answer (R):** During training, *Dropout* **randomly deactivates** a percentage of neurons in the hidden layers. This forces the network to avoid relying on any specific neuron or feature, improving its **generalization ability**.

ID: RN-21 | Topic: RN Types | Question (P): What is a **Convolutional Neural Network (CNN)** and what is it used for? | **Answer (R):** It is a network specialized for processing grid-like data (primarily **images**). It uses **convolutional layers** and *pooling* to automatically extract spatial features from the data.

ID: RN-22 | Topic: RN Types | Question (P): What is a **Recurrent Neural Network (RNN)** and what is its main characteristic? | **Answer (R):** It is a network designed to process **sequence data** (text, audio, time series). Its main characteristic is the **recurrent connection** (loop) that allows information from previous time steps to persist and be used in the current time step.

ID: RN-23 | Topic: Problems | Question (P): What is the **Vanishing Gradient** problem? | **Answer (R):** It occurs in deep networks when the gradient becomes progressively **smaller** as it backpropagates, causing the weights in the initial layers to barely update, and the network to stop learning effectively.

ID: RN-24 | Topic: Problems | Question (P): What is the **Exploding Gradient** problem? | **Answer (R):** It's the opposite of vanishing gradient. It occurs when gradients become **too large** during *Backpropagation*, leading to massive and unstable weight updates, causing the model to diverge or produce NaN (Not a Number).

ID: RN-25 | Topic: Problems | Question (P): What is a common solution for the Exploding Gradient problem? | **Answer (R):** A common technique is **gradient clipping**, where

gradients are limited to a predefined threshold if they exceed a certain value, preventing weight updates from becoming too large and unstable.